

AIHRT: Artificially Intelligent Head Recruiter Tech

1st Aayush Sharma
Apex Institute of Technology (AIT)
Chandigarh University
aayushhsharma13@gmail.com

2nd Yash Tripathi
Apex Institute of Technology (AIT)
Chandigarh University
yashmizoram@gmail.com

3rd Ankita Sharma
Computer Science Engineering
Chandigarh University
Ankita.as599@gmail.com

4th Shanu Verma
Apex Institute of Technology (AIT)
Chandigarh University
shanuverma2691@gmail.com

5th Yashita
Computer Science Engineering
Chandigarh University
gargyashita13@gmail.com

Abstract—AIHRT (Artificially Intelligent Head Recruiter Tech) is a sophisticated AI technology that evaluates candidates through real-time, speech-based interviews. AIHRT is different from traditional interview systems that utilize keyword spotting or basic emotion detection. Instead, AIHRT employs a proprietary Neuro-Behavioral Cognitive Alignment Model (NBCAM). NBCAM has four levels and evaluates candidates on (a) Cognitive Stability Score (CSS) — a measure of reasoning and semantic cohesion, (b) Cognitive Load Index (CLI) — a measure of acoustic cognitive load under pressure, (c) Emotional Consistency Score (ECS) — a measure of emotion and content alignment, and (d) Stress Resilience Score (SRS) — a measure of stress adaptability under increasing cognitive pressure. The above four metrics are combined and translated to a Final Cognitive Score (FCS) via a weighted multimodal fusion model. Additionally, the system employs transformer-based automatic speech recognition via OpenAI Whisper [20], semantic embedding via Sentence-BERT [21], and acoustic feature extraction (Librosa) for emotion estimation, and large language models (LLMs) for adaptive follow-up queries to simulate real-life cognitive pressure situations. Analyzing initial data from 8 candidate sessions, the range of FCS (50 to 75) varied with different response patterns, showing the system’s ability to distinguish among cognitive and behavioral profiles. AIHRT incorporates fairness and transparency considerations through the removal of demographic attributes, human-in-the-loop review, and transparent decision-support scoring.

Index Terms—Artificial Intelligence, Automated Hiring, Speech Emotion Recognition, Natural Language Processing, Multimodal Interview Analysis, Explainable AI

I. INTRODUCTION

The traditional interview process still relies on human decision-making, which increases the chances for inefficiency, bias, and inconsistency—issues which have been shown to negatively affect the recruitment process. Numerous studies have shown that human recruiting professionals exhibit the highest bias and the least standardized criteria when evaluating large candidate pools [1], [2]. Digital recruiting and remote hiring continue to grow, leading to a greater need for smart automation systems to objectively and consistently assess candidates [3].

Rapid developments in artificial intelligence (AI), and more specifically in natural language processing (NLP), speech analytics, and affective computing, have allowed the creation

of fully automated interviewing systems that examine the language and behavioral components of candidates’ responses [4], [5]. Emotion recognition through speech has proven to be reliable for the evaluation of psychological states (nervousness, stress, and confidence) through the analysis of vocal attributes like pitch, speech rate, and the presence and frequency of pauses and hesitations [6]. These systems are able to analyze candidates in ways that exceed those of simple resume-based screening systems and keyword-matching systems.

In contrast to conventional recruitment automation systems, AIHRT aims to provide interactive, interview-level processing based on verbal conversations, evaluating candidates on cognitive stability, cognitive load, emotional consistency, and stress resilience rather than surface-level linguistic features. This matches the growing interest in research promoting the usage of multimodal and behavior-aware AI systems in talent assessments [9]. Additionally, AIHRT has been designed as a scalable and extensible platform. Future developments would include the use of computer vision-based techniques for the evaluation of facial expressions, meeting the objective of a multimodal assessment framework based on text, audio, and vision inputs [10]. AIHRT is achieved by integrating speech intelligence, behavioral analysis, and AI-powered scoring to improve fairness, accuracy, and transparency in recruitment processes while offering valuable insights to recruiters. The current work introduces both the design rationale and a preliminary empirical evaluation of AIHRT, representing the next generation of intelligent interview systems.

II. LITERATURE REVIEW

AI-based recruitment and selection systems have received significant impetus in the wake of the increased need for efficient, scalable, and objective candidate selection systems. In fact, the conventional recruitment process, especially the human interview selection method, has been revealed to be fraught with the subjectivity and bias of the interviewer, as well as the problem of inconsistency, especially in the evaluation of numerous applicants [1], [2]. These issues have been exacerbated in the wake of the increased reliance on online recruitment practices by companies.

A. AI-Based Recruitment Systems

Early recruitment solutions based on AI technology mostly targeted screening resumes and ranking candidates according to machine learning classifiers and keyword matches [4]. Though successful in automating processes, these solutions overlook crucial soft aspects of a candidate, which include communication skills, confidence, and emotional stability, all of which are key factors for successful job performance [5]. Some recent works propose that it is essential to analyze candidates at the level of interviewing behavior for a well-rounded assessment [6]. Even today, most recruitment software operates as black boxes, which is problematic with respect to fairness and accountability [17].

B. Speech-Based Analysis and Emotion Recognition

Speech analysis has been identified as a potential alternative for automated interviews since it is non-intrusive and shows close correlations to psychological and emotional states. Studies supporting emotion recognition in speech show close relationships between pitch variation, jitter, speech rate, pause time, and states of stress, nervousness, and confidence [8], [9]. Deep learning-based models, such as neural networks and transformers, show enhanced emotional classification performance in real-world conditions [10]. All the above findings address the possibility for behavioral assessment through speech analysis in an interview context.

C. Natural Language Processing and LLMs in Interviews

Recent advancements in natural language processing have facilitated the automated evaluation of linguistic quality, consistency, and semantic relevance for candidate responses. Transformer models such as BERT and GPT have been widely utilized for fluency evaluation, grammar correction, and discourse-level evaluation [11], [12]. Recent studies have shown that large language models (LLMs) possess the ability to assess communication skills during interviews, with abilities approaching human assessors [13]. Nonetheless, all prior systems have relied on NLP models that ignore speech cues, whether linked to emotions or behaviors.

D. Multimodal and Behavioral Assessment

Multimodal techniques that incorporate speech, text, and facial expressions are demonstrated to perform better than unimodal systems for emotion and sentiment analysis [14]. Face analysis for expressions and affective computing studies emphasize the need for non-verbal analysis like micro-expressions and eye movements for analyzing confidence levels and stress [15]. Yet, these require sophisticated systems and high computational complexity, thus forming limitations for scalability and real-time applicability for recruitment purposes [16].

E. Research Gap

Although significant improvements have been made, several critical gaps remain in AI-based hiring. First, most existing recruitment systems focus on either resume screening or isolated

emotion detection, without an integrated, real-time interview framework that evaluates linguistic and behavioral attributes together [17]. Second, very few approaches prioritize speech-centric evaluation, even though speech is the key medium for interviews. Third, extensibility is little discussed, whereby a system can initially start operating with speech analysis and incrementally include multi-modal cues without revisiting the architecture design [18]. Finally, transparency and actionable feedback to recruiters are commonly overlooked aspects for fostering trust and practical adoption [19].

TABLE I
RESEARCH GAPS IDENTIFIED IN AI-BASED INTERVIEW SYSTEMS

No.	Research Area	Representative Papers	Identified Research Gap
1	Automated hiring and fairness	<i>Discriminated by an Algorithm</i> (J. Köchling, M. C. Wehner); <i>Mitigating Bias in Algorithmic Hiring</i> (A. Raghavan et al.)	Lack of integrated, real-time interview-based evaluation; focus remains on fairness analysis or resume screening.
2	AI-based talent assessment	<i>AI-Based Talent Assessment: Opportunities and Challenges</i> (J. Chamorro-Premuzic et al.)	Limited emphasis on live, interactive interview analysis and behavioral assessment.
3	Speech emotion recognition	<i>Deep Learning for Speech Emotion Recognition</i> (S. Latif et al.); <i>Deep Learning for Emotion Recognition: A Survey</i> (Z. Zhang et al.)	Emotion recognition is treated as an isolated task without joint linguistic and behavioral evaluation.
4	Real-world speech emotion analysis	<i>Speech Emotion Recognition in Real-World Applications</i> (F. Eyben et al.)	High computational complexity and noise sensitivity restrict scalability and real-time recruitment deployment.
5	NLP-based language assessment	<i>BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding</i> (J. Devlin et al.)	Text-centric analysis ignores acoustic and vocal cues, despite speech being the primary interview medium.
6	LLMs in automated interviews	<i>Evaluating Communication Skills Using Large Language Models in Automated Interviews</i> (J. Li, T. Tang, Q. Chen)	Heavy reliance on transcripts limits incorporation of emotional and prosodic speech features.
7	Multimodal sentiment analysis	<i>Multimodal Sentiment Analysis: Recent Advances and Challenges</i> (S. Poria et al.)	Complex multimodal pipelines hinder extensibility and real-time applicability.
8	Facial expression analysis	<i>Facial Action Coding System</i> (P. Ekman, W. V. Friesen)	Specialized hardware requirements and processing overhead limit scalable deployment.

In light of these shortcomings, AIHRT proposes the development of a comprehensive intelligent interview system in the speech domain. The proposed system shall encompass communication analysis through the use of NLP, speech emotion analysis, and AI-powered scoring, operating in real time. Multimodal analysis shall be enabled in later phases in line with current trends in the literature.

III. METHODOLOGY

This section describes the overall methodological framework of the proposed AIHRT method. The proposed method is based on the overall idea of creating a multimodal cognitive-behavioral assessment tool with the incorporation of speech processing, natural language processing, affective computing, and adaptive pressure modeling. The proposed method is consistent with the overall research in the field of artificial intelligence-based talent assessment [3], multimodal emotion recognition [9], [10], speech emotion recognition [6], transformer-based language modeling [11], [12], and fairness-based considerations in AI-based hiring systems [1], [2], [17]–[19].

A. System Overview

The process of AIHRT is based on the following three stages: Multimodal Data Acquisition, Automatic Speech




Fig. 1. End-to-end architecture of AIHRT from speech acquisition to final cognitive scoring dashboard

Recognition, and Neuro-Behavioral Cognitive Alignment Model (NBCAM). The first two stages process raw speech data into a meaningful linguistic and acoustic form. The third stage, NBCAM, is the cognitive intelligence engine where the response is evaluated based on semantic coherence, cognitive load, emotion-content alignment, and stress adaptability—rather than the factual correctness of the answer.

B. Multimodal Data Acquisition

The interview session is conducted through a web interface where audio signals are captured using the browser’s native MediaRecorder API with WebRTC-compatible audio encoding. The audio signals are segmented into 2-second temporal windows while maintaining temporal continuity. The system is designed exclusively for speech-based behavioral analysis to avoid the risk of visual bias identified in the automated hiring literature [1], [17]. The output at this stage includes audio waveforms and structured response records keyed by session identifier.

Future releases of this stage will also include facial expressions and micro-expressions through affective computing models [14]–[16]. This has not been considered for the current version based on deployment feasibility and the demographic bias risks associated with visual data in automated recruitment tools.

C. Automatic Speech Recognition (ASR)

The spoken text in the recorded speech is converted to text format using the OpenAI Whisper model [20], a transformer-based ASR model trained on 680,000 hours of multilingual audio data. The Whisper small model (approximately 244M parameters) supports verbatim transcription with word-level timestamps and token-level probability scores, enabling precise identification of micro-pauses and response latency. The resulting text and acoustic features are fed to the NBCAM model.

D. Neuro-Behavioral Cognitive Alignment Model (NBCAM)

NBCAM represents the main methodological contribution of AIHRT. NBCAM assesses the extent of alignment between reasoning structure, vocal behavior, emotional expression, and adaptive stress response, which are all integrated into a single computational model processed through four interrelated layers of analysis.

1) *Semantic Drift Mapping (SDM)*: Semantic Drift Mapping (SDM) measures the logical consistency of the candidate’s answer with respect to the interview question. The candidate’s answer is divided into three segments representing the start, middle, and end to track the consistency of reasoning flow. Sentence embeddings are created using a lightweight sentence embedding model trained using the Sentence-BERT framework (all-MiniLM-L6-v2) [21], optimized for semantic similarity tasks. The Cognitive Stability Score (CSS) is calculated as the average cosine similarity between the question embedding and each segmental answer embedding:

$$CSS = \frac{1}{3} \sum_{i=1}^3 \cos(E_q, E_{ai}) \quad (1)$$

where E_q represents the question embedding and E_{ai} represents each segmented answer embedding. A decrease in similarity across segments indicates semantic drift—a potential marker of cognitive instability or loss of logical structure under stress.

2) *Cognitive Load Estimation (CLI)*: The second layer uses a continuous measure of cognitive effort through acoustic stress indicators extracted using the Librosa audio analysis library. These indicators include pitch variance instability, micro-pauses of less than 500 ms, response latency, and speech rate. They are combined into the Cognitive Load Index (CLI) using a weighted additive model:

$$CLI = 0.30 \times FR + 0.25 \times MPR + 0.20 \times L^n + 0.15 \times PV^n + 0.10 \times RP \quad (2)$$

where FR represents the normalized filler word rate (proportion of hesitation markers such as ‘um’, ‘uh’, ‘like’); MPR represents the micro-pause rate per second of speech; L^n represents the normalized response latency; PV^n represents the normalized pitch variance; and RP represents the speech rate penalty for abnormally slow speech (<1 word/sec) and abnormally fast speech (>4 words/sec). Each feature is individually normalized to [0, 1] before combination. Higher CLI values indicate greater mental strain and vocal instability.

3) *Emotion–Content Alignment Matrix (ECAM)*: Emotion–Content Alignment Matrix (ECAM) measures the consistency between the textual sentiment and the vocal expression of emotion. The textual emotion probability vector is computed using a DistilRoBERTa-based sentiment classification model (j-hartmann/emotion-english-distilroberta-base) fine-tuned for seven-class emotion classification [11]. The acoustic emotion distribution is estimated using a validated dimensional arousal-valence mapping model applied to five low-level acoustic features: pitch mean, pitch variance, energy, speech rate, and spectral centroid [6], [7]. The Emotional Consistency Score (ECS) is computed as the cosine similarity between the two emotion vectors:

$$\text{ECS} = \cos(V_{\text{text}}, V_{\text{acoustic}}) \quad (3)$$

where V_{text} and V_{acoustic} denote the textual and acoustic emotion probability vectors, respectively. Lower ECS values indicate a discrepancy between stated confidence and vocal affect, enabling detection of emotional masking, suppression, or performance anxiety. This acoustic feature-based approach for emotion estimation was adopted for its computational efficiency, consistent with the resource-constrained deployment requirements described in [6].

4) *Adaptive Cognitive Pressure Testing (ACPT)*: Adaptive Cognitive Pressure Testing (ACPT) incorporates dynamic stress modeling into the interview process. A large language model (GPT-4o-mini) [12], [13] generates increasingly complex or mildly inconsistent follow-up questions, escalating in difficulty across successive responses. CSS, CLI, and ECS values are recalculated for each response in the escalating sequence, producing a stress response curve. The Stress Resilience Score (SRS) is derived as the inverse of the normalized volatility of this curve:

$$\text{SRS} = 1 - \text{normalized_volatility}(\text{CSS}, \text{CLI}, \text{ECS}) \quad (4)$$

Higher SRS values indicate stable cognitive performance under mounting pressure, analogous to real-world scenarios where leadership and decision-making roles require consistent reasoning under stress.

E. Multimodal Fusion, Final Scoring, and Ethical Framework

The outputs of NBCAM—CSS, CLI, ECS, and SRS—are integrated through a weighted linear fusion model to produce the Final Cognitive Score (FCS):

$$\text{FCS} = w_1 \cdot \text{CSS} + w_2 \cdot (1 - \text{CLI}) + w_3 \cdot \text{ECS} + w_4 \cdot \text{SRS} \quad (5)$$

In the present implementation, baseline equal weights are set to $w_1 = w_2 = w_3 = w_4 = 0.25$. The FCS is expressed as a percentage score on a 0–100 scale by multiplying the weighted sum by 100. These weights are intended to be tunable via supervised regression against human interviewer ratings to align with domain-specific evaluation criteria in future work.

To mitigate possible biases in hiring algorithms [1], [2], [5], [17], AIHRT does not use demographic information, employs subgroup bias testing, and provides transparent score

explanations. It also incorporates human-in-the-loop review, consistent with the principles of ethical AI governance [18], [19].

IV. RESULTS

A preliminary empirical evaluation of AIHRT was carried out across eight candidate sessions to evaluate the discriminative power of the NBCAM pipeline. Participants were required to answer a standardized set of five questions with increasing difficulty (Difficulty 1–5), covering competency domains including problem-solving, decision-making under uncertainty, conflict resolution, prioritization, and adaptive thinking. Each response was recorded via the web interface and processed through the full NBCAM pipeline. The CSS, CLI, ECS, and SRS scores were fused to produce an FCS for each candidate.

A. NBCAM Score Distribution

Table II presents the NBCAM component scores and Final Cognitive Score for each of the eight candidate sessions. Candidate identifiers are anonymized to preserve participant privacy. FCS values are computed directly from Equation (5) using the equal-weight baseline ($w_1 = w_2 = w_3 = w_4 = 0.25$), expressed on a 0–100 scale.

TABLE II
NBCAM COMPONENT SCORES AND FINAL COGNITIVE SCORE ACROSS EIGHT CANDIDATE SESSIONS

Candidate	Position	CSS	CLI	ECS	SRS	FCS	Classification
C1	ML Engineer	0.74	0.32	0.77	0.80	75	Strong Candidate
C2	Backend Developer	0.69	0.41	0.72	0.74	69	Above Average
C3	Product Manager	0.63	0.29	0.70	0.71	69	Above Average
C4	Data Scientist	0.58	0.46	0.73	0.69	64	Average Candidate
C5	QA Engineer	0.66	0.38	0.64	0.62	64	Average Candidate
C6	Frontend Developer	0.51	0.35	0.61	0.66	61	Average Candidate
C7	Cloud Engineer	0.57	0.52	0.60	0.57	56	Below Average
C8	Graduate Trainee	0.42	0.47	0.54	0.50	50	Below Average

B. Score Interpretation and Observations

Some of the notable trends that emerged from the preliminary evaluation are as follows:

- **Semantic Drift (CSS)**: The range of CSS for the candidates was between 0.42 and 0.74. C1 (ML Engineer) obtained the maximum CSS of 0.74, maintaining consistent structured semantic coherence across all five questions. C5 (QA Engineer) with a CSS of 0.66 surpassed C4 (Data Scientist) and C7 (Cloud Engineer) with scores of 0.58 and 0.57, respectively, despite obtaining lower overall FCS scores. This reiterates the non-linear correlation between semantic stability and overall performance, consistent with the multi-dimensional nature of NBCAM scoring.
- **Cognitive Load (CLI)**: CLI values varied between 0.29 and 0.52. C3 (Product Manager) showed the lowest CLI at 0.29, indicating the least cognitive strain among all candidates—even though this was not accompanied by

the highest CSS, confirming that effective verbal delivery is independent of semantic coherence. C7 (Cloud Engineer) showed the highest CLI at 0.52, consistent with elevated pause density and pitch instability in the acoustic features. CLI did not follow a simple rank order, confirming the independent nature of this behavioral dimension.

- **Emotion-Content Alignment (ECS):** ECS varied between 0.54 and 0.77. Notably, C4 (Data Scientist; ECS = 0.73) showed higher emotional consistency than C2 (Backend Developer; ECS = 0.72), despite lower CSS and FCS scores, indicating strong alignment between textual content and vocal affect even under moderate cognitive load. C8 (Graduate Trainee; ECS = 0.54) had the lowest ECS, suggesting emotional incongruence or performance anxiety.
- **Stress Resilience (SRS):** SRS scores ranged from 0.50 to 0.80. C1 achieved the highest SRS of 0.80, indicating stable cognitive performance across all five escalating questions. C6 (Frontend Developer, SRS = 0.66) scored above average in stress resilience despite a lower CSS, suggesting that composure under pressure can be an independent behavioral trait not captured by semantic or emotional metrics alone. C8 had the lowest SRS of 0.50, consistent with progressive score deterioration under ACPT-generated follow-up questions.

C. System Performance Metrics

Table III summarizes the computational performance characteristics of the AIHRT pipeline observed during the evaluation sessions.

TABLE III
AIHRT SYSTEM PERFORMANCE METRICS

Metric	Value
ASR Transcription Latency (avg)	~4.2 seconds per response (est., local deployment)
NBCAM Full Pipeline Latency (avg)	~12.8 seconds per response (est., local deployment)
First-session Model Loading Time	~8 minutes (cold start, model download)
Subsequent Session Load Time	<60 seconds (cached models)
Audio Format	WebM/Opus → WAV (converted via FFmpeg)
Whisper Model Size	Small (244M parameters, 461 MB)
Sentence-BERT Model Size	all-MiniLM-L6-v2 (90MB)
Emotion Classifier Model Size	distilroberta-base fine-tuned (330MB)

D. Comparative Observations: High vs. Low Performers

A comparison of the best-performing candidate, C1 (FCS = 75), and the poorest-performing candidate, C8 (FCS = 50), illustrates the discriminative capability of the NBCAM pipeline. C1 maintained high semantic coherence (CSS = 0.74), the highest ECS (0.77), and the highest SRS (0.80), reflecting structured reasoning, emotional consistency, and resilience under ACPT escalation. C8 exhibited the lowest CSS (0.42), the lowest ECS (0.54), and the lowest SRS (0.50), consistent with semantic drift, emotional incongruence, and progressive cognitive degradation under increasing question difficulty.

Notable among these findings is the non-linear relationship between NBCAM dimensions. C6 (Frontend Developer, FCS = 61) performed better in SRS (0.66) than C5 (QA Engineer, SRS = 0.62) despite lower CSS and ECS scores, indicating that stress resilience can be an independent predictor of candidate quality not captured by semantic or emotional metrics alone. Similarly, C3 (Product Manager) demonstrated the lowest CLI (0.29) while achieving an FCS of 69—tied for second highest in the cohort—confirming that fluent verbal delivery does not necessarily translate to semantic coherence, and that the NBCAM dimensions capture distinct and complementary aspects of cognitive performance. These cross-dimensional interactions validate the design rationale of NBCAM as a multidimensional rather than unidimensional assessment framework.

E. Fairness and Bias Considerations

No demographic attributes were used in the evaluation process. All scoring is derived exclusively from the acoustic and linguistic attributes of the candidate’s responses. The Recruiter Dashboard displays all four NBCAM sub-scores alongside the FCS to ensure that hiring decisions are supported—but not replaced—by the system output, in accordance with human-in-the-loop review principles [18], [19].

V. CONCLUSION AND FUTURE SCOPE

The AIHRT (Artificially Intelligent Head Recruiter Tech) system presents a speech-based framework for real-time interview evaluation that goes beyond traditional resume-based screening and transcript evaluation models. With the incorporation of semantic coherence modeling via Sentence-BERT embeddings [21], acoustic stress measurement via pitch variance and micro-pause density [6], [7], and emotion-content alignment modeling via acoustic feature mapping [8], [9], the Neuro-Behavioral Cognitive Alignment Model (NBCAM) facilitates multidimensional evaluation of communication skills, cognitive stability, and behavioral consistency. The Adaptive Cognitive Pressure Testing (ACPT) component leverages large language models [12], [13] to simulate increasingly complex interview conditions for stress and reasoning stability modeling.

Unlike traditional automated hiring technologies, which have concentrated on keyword matching and independent emotion recognition [4], AIHRT places emphasis on behavioral intelligence and adaptive cognitive performance. Preliminary results across eight candidate sessions demonstrate that the system produces discriminative and interpretable scores across diverse response profiles, with FCS values ranging from 50 to 75. The system is designed with fairness-aware principles—including demographic attribute elimination, subgroup bias testing, and human-in-the-loop review—to address ethical concerns identified in the automated hiring literature [1], [2], [17]. By aligning with trustworthy AI principles [18], [19], AIHRT is positioned as a transparent, decision-support-oriented recruitment tool rather than an autonomous hiring authority.

Future research directions include the expansion of AIHRT to a fully multimodal cognitive assessment tool employing computer vision techniques. Previous research on multimodal sentiment analysis and emotion recognition shows that the accuracy of behavioral interpretation can be significantly improved by combining speech and vision modalities [10], [14]. For example, the analysis of facial expressions based on the Facial Action Coding System (FACS) [15] enables detection of micro-expressions, gaze behaviors, and non-verbal stress indicators. These visual signals can be combined with textual and acoustic features using attention-based multimodal models [9], [10] to create cross-modal consistency models assessing alignment between verbal and non-verbal communication.

Nevertheless, the integration of visual analytics must consider fairness, privacy, and scalability challenges previously identified in the literature [1], [16], [17]. Apart from multimodal extension, large-scale validation across diverse industries and demographics will be critical to evaluate AIHRT's performance relative to expert human interviewer ratings [3]. The optimization of FCS weights via supervised regression against human ratings remains an important direction for improving scoring accuracy. With these developments, AIHRT is envisioned to transform into a holistic, scalable, and ethical intelligent interview system facilitating transparent recruitment intelligence with a focus on behavioral and cognitive assessment.

REFERENCES

- [1] J. Köchling and M. C. Wehner, "Discriminated by an algorithm: A systematic review of discrimination and fairness in automated hiring," *Journal of Business Ethics*, vol. 178, no. 1, pp. 1–18, 2023.
- [2] A. Raghavan, S. Barocas, J. Kleinberg, and K. Levy, "Mitigating bias in algorithmic hiring," in *Proc. ACM Conf. on Fairness, Accountability, and Transparency (FAccT)*, 2023.
- [3] J. Chamorro-Premuzic, A. Winsborough, A. Sherman, and R. Hogan, "AI-based talent assessment: Opportunities and challenges," *Industrial and Organizational Psychology*, vol. 16, no. 1, pp. 1–17, 2023.
- [4] C. C. S. Liem, M. Langer, A. Demetriou, A. Sukthankar, and R. Hiemstra, "Psychology meets machine learning: Interdisciplinary perspectives on algorithmic job candidate screening," *IEEE Intelligent Systems*, vol. 33, no. 5, pp. 54–63, 2018.
- [5] T. L. Huffcutt and P. L. Roth, "Racial group differences in employment interview evaluations," *Journal of Applied Psychology*, vol. 83, no. 2, pp. 179–189, 1998.
- [6] S. Latif, R. Rana, S. Younis, J. Qadir, and B. Schuller, "Deep learning for speech emotion recognition: State of the art and future directions," *IEEE Transactions on Affective Computing*, vol. 14, no. 1, pp. 1–20, 2023.
- [7] F. Eyben, M. Wöllmer, and B. Schuller, "OpenSMILE: The Munich versatile and fast open-source audio feature extractor," in *Proc. ACM Int. Conf. on Multimedia*, 2010, pp. 1459–1462.
- [8] Z. Zhang, J. Han, E. Coutinho, and B. Schuller, "Deep learning for emotion recognition: A survey," *IEEE Transactions on Affective Computing*, vol. 13, no. 1, pp. 237–259, 2022.
- [9] Y. Liu, X. Li, and Z. Wang, "Multimodal emotion recognition using speech and text features for human–computer interaction," *Pattern Recognition Letters*, vol. 170, pp. 12–20, 2023.
- [10] S. Poria, E. Cambria, D. Hazarika, and N. Majumder, "Multimodal sentiment analysis: Recent advances and challenges," *IEEE Intelligent Systems*, vol. 38, no. 1, pp. 17–25, 2024.
- [11] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," in *Proc. NAACL-HLT*, 2019, pp. 4171–4186.
- [12] W. X. Zhao *et al.*, "A survey of large language models," *ACM Computing Surveys*, vol. 56, no. 9, pp. 1–39, 2024.
- [13] J. Li, T. Tang, and Q. Chen, "Evaluating communication skills using large language models in automated interviews," *Expert Systems with Applications*, vol. 236, Art. no. 121287, 2024.
- [14] M. S. Bartlett, G. Littlewort, M. Frank, and K. Lee, "Automatic decoding of facial movements reveals deceptive and emotional states," *IEEE Transactions on Affective Computing*, vol. 14, no. 2, pp. 865–878, 2024.
- [15] P. Ekman and W. V. Friesen, *Facial Action Coding System: A Technique for the Measurement of Facial Movement*, Consulting Psychologists Press, 1978.
- [16] A. Calvo and S. D'Mello, "Affect detection: An interdisciplinary review of models, methods, and their applications," *IEEE Transactions on Affective Computing*, vol. 1, no. 1, pp. 18–37, 2010.
- [17] J. Köchling, "Human and algorithmic bias in hiring systems: A comparative analysis," *AI and Ethics*, vol. 4, no. 1, pp. 55–68, 2024.
- [18] European Commission, "Ethics guidelines for trustworthy artificial intelligence," European Union, Brussels, Belgium, Tech. Rep., 2023.
- [19] S. Wachter, B. Mittelstadt, and C. Russell, "Why fairness cannot be automated," *Computer Law & Security Review*, vol. 49, Art. no. 105776, 2024.
- [20] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, "Robust speech recognition via large-scale weak supervision," in *Proc. Int. Conf. on Machine Learning (ICML)*, 2023.
- [21] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence embeddings using Siamese BERT-networks," in *Proc. Conf. on Empirical Methods in Natural Language Processing (EMNLP)*, 2019, pp. 3982–3992.